

# Lab Exercises Week 02

Alison Lawyer

2024-09-23

```
knitr::opts_chunk$set(echo = TRUE)
knitr::opts_chunk$set(tidy.opts = list(width.cutoff = 60), tidy = TRUE)
```

## LAB EXERCISES

This week you will continue practicing basic R syntax, functions, and descriptive summary statistics. You will also make a pretty table that you could use to report your results in a document or presentation. For all exercises, we will use the palmer penguins data that you started using for homework 1.

### Part 1

First, let's look at NA values a bit more closely

```
library(tidyverse)
```

```
## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr      1.1.4      v readr      2.1.5
## v forcats    1.0.0      v stringr   1.5.1
## v ggplot2    3.5.1      v tibble    3.2.1
## v lubridate  1.9.3      v tidyr     1.3.1
## v purrr      1.0.2
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
library(dplyr)
```

```
# Load the palmer penguin dataset
```

```
library(palmerpenguins)
```

```
penguins <- data.frame(penguins)
```

```
# You can use the function is.na() to test whether any
```

```
# value in a vector of values is NA Check for NAs in the
```

```
# column 'body_mass_g'
```

```
body_mass_na <- is.na(penguins$body_mass_g)
```

```
print(body_mass_na)
```

```
## [1] FALSE FALSE FALSE TRUE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [13] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [25] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [37] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [49] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [61] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [73] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [85] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [97] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [109] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [121] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [133] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [145] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [157] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [169] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [181] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [193] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [205] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [217] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [229] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [241] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [253] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [265] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE TRUE FALSE FALSE FALSE
## [277] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [289] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [301] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [313] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [325] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [337] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
```

```
# It will be difficult to detect the TRUE values in a sea
# of FALSE values Since TRUE/FALSE are internally stored as
# 1/0 values, we can do numerical calculations on the
# result of is.na() Let's try that with sum() function by
# wrapping the is.na() function in a sum() function
# Wrapping a function means using one function inside of
# another function
body_mass_na_sum <- sum(is.na(penguins$body_mass_g))
print(body_mass_na_sum)
```

```
## [1] 2
```

```
# This is the number of rows with missing values for body
# mass. We can use this to incorporate the calculation in
# a summarize() function too. Create a summary by species
# using the summarize() function, where a new variable
# holds the number of missing values for 'body_mass_g'.
grouped_missing_body_mass <- penguins %>%
  group_by(species) %>%
  summarize(sum_missing_body_mass = sum(is.na(body_mass_g)))
head(grouped_missing_body_mass)
```

```
## # A tibble: 3 x 2
```

```
## species    sum_missing_body_mass
## <fct>      <int>
## 1 Adelie      1
## 2 Chinstrap   0
## 3 Gentoo      1
```

*# Now let's add to that also the proportion of missing values for each species, that is, the count of missing values divided by the total number of rows for each species. Add that to your code above and re-create the summary with two variables, one for number of missing rows and one for percentage of missing rows. Both variables will be created in the same summarize() function statement. Hint: you can use the function n() to count the number of rows.*

```
grouped_proportion <- penguins %>%
  group_by(species) %>%
  summarize(total_missing_rows = sum(is.na(body_mass_g)), proportion_missing = (sum(is.na(body_mass_g)) /
    100)
head(grouped_proportion)
```

```
## # A tibble: 3 x 3
## species    total_missing_rows proportion_missing
## <fct>      <int>          <dbl>
## 1 Adelie      1          0.658
## 2 Chinstrap   0          0
## 3 Gentoo      1          0.806
```

*# When dealing with multiple variables that each may have missing values, one way to filter the dataset to only those rows with no missing values in any variable is to use the function complete.cases(). Look up the help for this function by typing ?complete.cases into your console. See if you can figure out how to use this function. If it is not obvious from the help page, use other resources online or play around until you get it to work (this is a very common way to solve coding problems!).*

```
penguins_without_missing_rows <- penguins[complete.cases(penguins),
]
# penguins_without_missing_rows <- penguins %>%
# filter(complete.cases(penguins))
# penguins_without_missing_rows <- na.exclude(penguins)
# penguins_without_missing_rows <- na.omit(penguins)
head(penguins_without_missing_rows)
```

```
## species    island bill_length_mm bill_depth_mm flipper_length_mm body_mass_g
## 1 Adelie Torgersen    39.1         18.7          181          3750
## 2 Adelie Torgersen    39.5         17.4          186          3800
## 3 Adelie Torgersen    40.3         18.0          195          3250
## 5 Adelie Torgersen    36.7         19.3          193          3450
## 6 Adelie Torgersen    39.3         20.6          190          3650
## 7 Adelie Torgersen    38.9         17.8          181          3625
```

```
##      sex year
## 1   male 2007
## 2 female 2007
## 3 female 2007
## 5 female 2007
## 6   male 2007
## 7 female 2007
```

## Part 2

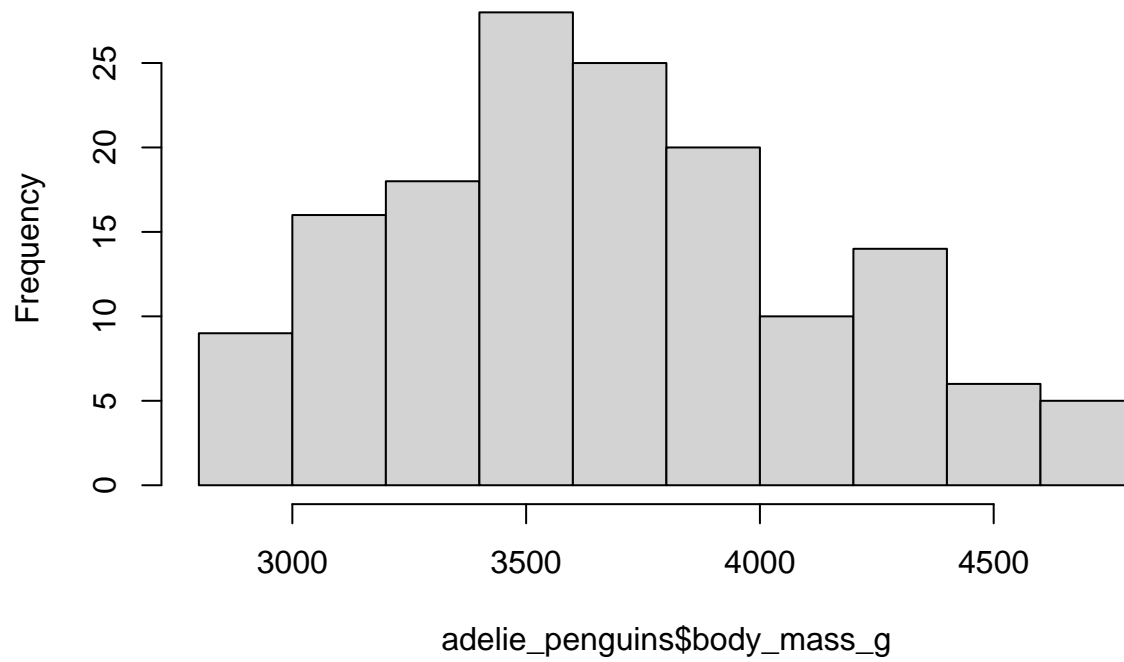
Now let's look at the distribution of body mass for each species by making histograms

```
# Create a new dataframe holding body mass data for Adelie
# penguins only. Keep the species, sex, and body_mass_g
# columns only
adelie_penguins <- penguins %>%
  filter(species == "Adelie") %>%
  select(c(species, sex, body_mass_g))
gentoo_penguins <- penguins %>%
  filter(species == "Gentoo") %>%
  select(c(species, sex, body_mass_g))
chinstrap_penguins <- penguins %>%
  filter(species == "Chinstrap") %>%
  select(c(species, sex, body_mass_g))
head(gentoo_penguins)
```

```
##   species    sex body_mass_g
## 1  Gentoo female      4500
## 2  Gentoo  male      5700
## 3  Gentoo female      4450
## 4  Gentoo  male      5700
## 5  Gentoo  male      5400
## 6  Gentoo female      4550
```


```
# Use the function hist() to make a histogram of weight for
# Adelie penguins
hist(adelie_penguins$body_mass_g)
```

**Histogram of adelie\_penguins\$body\_mass\_g**



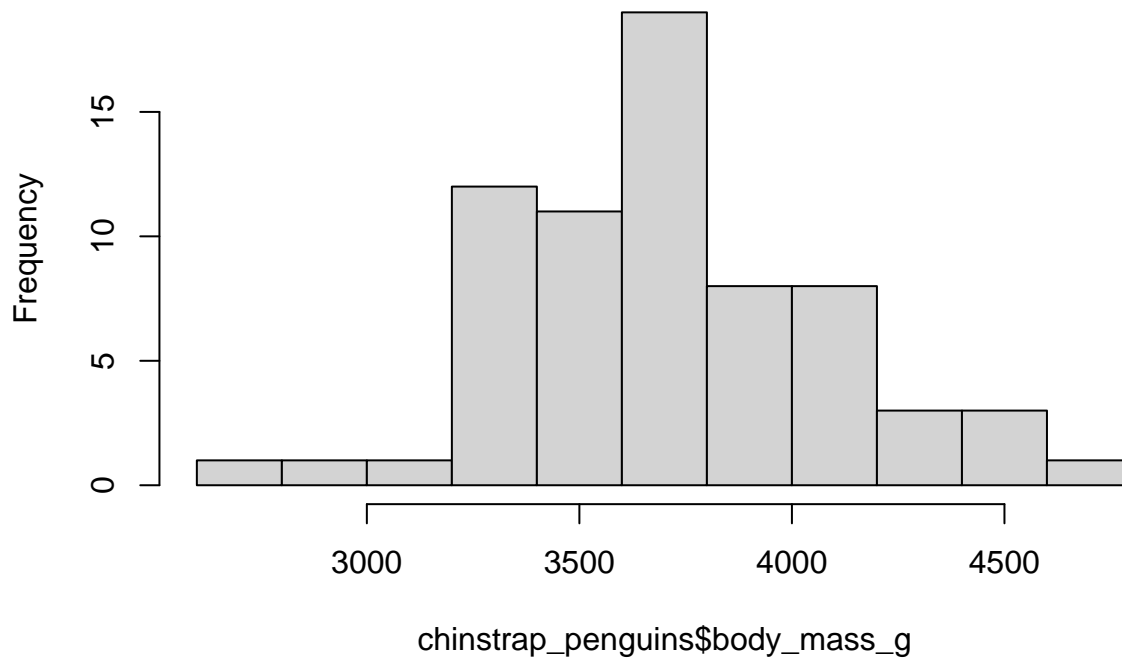
```
hist(gentoo_penguins$body_mass_g)
```

**Histogram of gentoo\_penguins\$body\_mass\_g**



```
hist(chinstrap_penguins$body_mass_g)
```

## Histogram of chinstrap\_penguins\$body\_mass\_g



```
# Is there skew in the data, and if so are the data left or  
# right skewed? plot mean and median lines to help  
# visualize
```

```
# Repeat this for the other two species and answer the same  
# question
```

### Part 3

Now, let's practice adding some descriptive statistics of body mass and save everything in a summary dataframe.

```
# Create a summary dataframe of the variable body_mass_g by  
# species that contains the following descriptive stats:  
# Mean, Sample Variance, Sample Standard Deviation, Median,  
# Interquartile Range (Tip: look up function IQR() to do  
# this), Min, Max, Number of observations
```

```
penguins %>%  
  group_by(species) %>%  
  summarize(mean = mean(body_mass_g, na.rm = TRUE), variance = var(body_mass_g,  
    na.rm = TRUE), sd = sd(body_mass_g, na.rm = TRUE), median = median(body_mass_g,  
    na.rm = TRUE), IQR = IQR(body_mass_g, na.rm = TRUE),
```

```
min = min(body_mass_g, na.rm = TRUE), max = max(body_mass_g,
na.rm = TRUE), count = n())
```

```
## # A tibble: 3 x 9
##   species    mean variance      sd median   IQR   min   max count
##   <fct>      <dbl>    <dbl> <dbl> <dbl> <dbl> <int> <int> <int>
## 1 Adelie    3701.    210283.  459.   3700   650   2850  4775   152
## 2 Chinstrap 3733.    147713.  384.   3700   462.   2700  4800    68
## 3 Gentoo    5076.    254133.  504.   5000   800   3950  6300   124
```

## Part 4

Finally, let's make a more automated data summary in form of a tidy table using the `table1` package

```
# Make a pretty table of your data summary with the
# table1() function The table should list species as
# columns and summary statistics for body mass as rows
library(table1)
```

```
##
## Attaching package: 'table1'

## The following objects are masked from 'package:base':
##
##   units, units<-
```

```
table1(~body_mass_g | species, data = na.omit(penguins))
```

## Get nicer 'table1' LaTeX output by simply installing the 'kableExtra' package

	Adelie	Chinstrap	Gentoo	Overall
	(N=146)	(N=68)	(N=119)	(N=333)
body__mass__g				
Mean (SD)	3710 (459)	3730 (384)	5090 (501)	4210 (805)
Median [Min, Max]	3700 [2850, 4780]	3700 [2700, 4800]	5050 [3950, 6300]	4050 [2700, 6300]